

Assignment Project Exam Help

Algorithms
Vectorization

<https://eduassistpro.github.io>

BTRY/STSCI 4

Add WeChat edu_assist_pr

Assignment Project Exam Help

- One of the key concerns of computer science – how long does

- <https://eduassistpro.github.io>

- Relevant to data size: I can produce an answer for a data set of size 100, but how long will data of size 100,00

- Usually calculated in terms of the number of operations required.

Example: The sort Problem

One of the classic problems in computer science.

- We have n numbers x_1, \dots, x_n .
- We want to put them in order so that

- <https://eduassistpro.github.io>
- How does computing time change as n increases?

Multiple ways to do this:

- Selection Sort
- Insertion Sort
- Bubble Sort
- Quick Sort

and others – how you do this makes a difference!

A First Problem – Finding the Minimum

Suppose that we just want $\min(x_1, \dots, x_n)$.

Program 1: loop through and check if each x_j is the minimum.

```
findmin = function(x){  
  foundmin = FALSE # Have we found the minimum?  
  i = 0  
  while(  
    ismin = TRUE # Assume x[i] is the minimum  
    for(j in 1:length(x)){ # Check against all ot  
      if(x[j] < x[i]){ ismin = FALSE }  
    }  
    # If nothing is less than x[i] it must be what we wan  
    if(ismin){ foundmin=TRUE }  
  }  
  return(x[i])  
}
```

An Analysis of Computing Time

We will count the number of comparisons made.

- At each i , we compare $x[i]$ to n other entries.

-

- <https://eduassistpro.github.io>

- What about somewhere in the middle?

- If we get about half way, we consider $n^2/2$ comparisons.

- If we think about x being random
have to look at a number of entries *proportional* to n .

A bit complicated; how do we simplify this?

Order Notation

- Suppose Algorithm 1 takes $3n^3 - 6n + 2$ operations and Algorithm 2 takes $4n^2 + 3$ operations.
- If n is large enough $3n^3 - 6n + 2 > 4n^2 + 3$, so Algorithm 1

- <https://eduassistpro.github.io> and +3.

- In fact, the 3 and 4 don't matter either. $an^3 + bn^2 + cn + d$ will always be dominated by an^3 if $a > 0$.
- We say that $an^3 + bn^2 + cn + d$ is $O(n^3)$.
- For the minimum search above, our algorithm requires $O(n^2)$ comparisons.

Assignment Project Exam Help

We will not worry a great deal about the formalism, but

- “Big-O” notation:

gh and

- <https://eduassistpro.github.io>

- We also have little-o notation (will come up later)

$f(x) = o(g(x))$ if $|f(x)/g(x)| \rightarrow 0$

- For example $1/n^2 = o(1/n)$.
- I.e., Big-O means “bounded by”, little-o means “much less than”.

General Rules

Most expressions in terms of x^α , e^x , $\log(x)$

- If $x \rightarrow \infty$, $\alpha_1 > \alpha_2 > 0$ then for x large enough

$$e^x > x^{\alpha_1} > x^{\alpha_2} > \log(x)$$

■ <https://eduassistpro.github.io>

■ does.

- If $x \rightarrow 0$ (looking at small numbers) then

$$|x^{\alpha_1}| < |$$

expression dominated by the smallest power of x .

Note $e^x \rightarrow 1$, $|\log(x)| \rightarrow \infty$; even larger than powers if these appear.

A More Efficient Search

$O(n^2)$ is pretty bad, can we make this better? Keep track of the

minimum *up-till-now*

```
FindMin2 = function(x){
```

```
  mi
```

```
  mi
```

```
  fo
```

```
    min = x[i] # Update minimum if x[i] is
```

```
    min.i = i # less than current value
```

```
  }
```

```
}
```

```
return(list(min=min,min.i=min.i))
```

```
}
```

Only does $n - 1 = O(n)$ comparisons.

Selection Sort

Now that we can find the minimum easily. Sort by continually finding the minimum:

```
SelectionSort = function(x){  
  y = 0*x    # Store the sorted vector  
  in  
  n = le  
  fo  
    cur.min = FindMin2(x) # Find and record the current  
    y[i] = cur.min$min    # minimum in x  
    ind[i] = cur.min$min.i  
    x = x[-cur.min$min.i] # Delete that ent  
  }  
  y[n] = x          # Fix last element.  
  return( list(y = y,ind = ind) )  
}
```

Analyzing Selection Sort

■ First iteration – it takes $n - 1$ entries to find the minimum in x .

■ Record these and remove them from x .

■ Next iteration, x now length $n - 1$, so we have $n - 2$

■ <https://eduassistpro.github.io>

$$\sum_{k=1}^n (n - k) = \frac{1}{2}n(n - 1)$$

Add WeChat edu_assist_pro

comparisons.

■ So if my data is 10 times as long, I have to put in 100 times the effort to sort it.

Insertion Sort

No R code this time:

- Start with x and assign y .
- Take each element of x in turn, insert it into y so y is sorted.
- After k steps:

<https://eduassistpro.github.io>

so that $y_j \leq x_1 \leq y_{j+1}$.

- Might not need to compare x_1 to y_j , stop at first such that $y_j > x_1$.
- Configuration of x changes number of comparison (what's fastest?)
- Tends to be faster than Selection Sort; but still generally $O(n^2)$.

Bubble Sort

Assignment Project Exam Help

It might be useful not to need to store a new vector; instead just swap entries

Repe

- <https://eduassistpro.github.io>
-

until x is sorted (and you make no more swaps).

Still $O(n^2)$ (exercise: what's the worst case?)

Tends to be slow; speed generally more an issue than memory.

Quick Sort

Due to Hoare (1960):

- Divide the data in two:

- 1 Those less than `x[1]`; call this `a`

- <https://eduassistpro.github.io>

- Split `a` and produce `c(d,a[1],`

- [Add WeChat edu_assist_pr](#)
- Eventually we have only one element – that's
- Nice Wikipedia animation.
- But how are we going to set this scheme up?

Graphically

Divide and conquer:

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Recursive Programming

It's ok to have a function call itself!

```
QuickSort = function(x){  
  if (length(x) == 1 || is.null(x)) return(x) }
```

```
  lo
```

```
  up
```

```
  for
```

```
    if(x[i] <= x[1]){ lower = c(lower,x[i]
```

```
    else{ upper = c(upper,x[i]) }  
  }
```

```
  lower = QuickSort(lower) # Now sort each of these
```

```
  upper = QuickSort(upper) # and put them back together
```

```
  return( c(lower,x[1],upper) )  
}
```


Graphically

Strategy goes left to right:

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Analyzing QuickSort

Suppose that we (luckily!) exactly partition the data set in 2 each time.

- At first level, I make $n - 1$ comparisons.
- At second level, I make two lots of $(n - 1)/2 - 1$ comparisons,

- <https://eduassistpro.github.io>

- So every level has $O(n)$ comparisons, but if there are $n = 2^k$ objects, there are $k = \log(n)/\log$

- That means the total cost is $O(n \log n)$ (much better!).

- Worse case: x already sorted, then we still have $O(n^2)$.
- Start by randomly permuting x : expected cost is still $O(n \log n)$.

To iterate is human, to recurse divine! - L. Peter Deutsch

Graphically

At each level, divide the data by 2, but twice as many nodes; but $\log_2(n)$ levels.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Why Should Statisticians Care?

n usually = size of data set

- Operations like sorting (and many others) are integral to parts of statistical computing.
- For an $O(n^2)$ operation, something feasible at $n = 100$ is not

■ <https://eduassistpro.github.io>

■ web-commerce, social networking, brain images and remote sensing, high-throughput astronomical surveys, citizen science,...

- Each now produces either millions of records, or hundreds of thousands of variables, or both.
- Historically: data sets grow as fast as computing speed.
- Lesson: if it isn't $O(n)$, in the long-run it will be too slow. (but note the long run can be some time away)

P and NP

- Much of the topic of *algorithms* in CS devoted to computational complexity.

- Not always easy to calculate.



<https://eduassistpro.github.io>

- polynomial time algorithms.

- Problems (eg how to sort a vector) divided into

P The set of problems that can be solved in polynomial time.

NP The set of problems for which a solution can be *verified* in polynomial time (eg: is this vector sorted?)

Example and a Question

Traveling salesman problem

- Salesman must visit all of a set of cities.

-

- <https://eduassistpro.github.io>

A solution is easy to check; but finding out if there is one is

Question:

Clearly anything in P is in NP , but what a
around?

One of the great unsolved problems of mathematics.

NP Hard

Formally, NP Hard is defined in terms of reducing NP problems to NP Hard problems (eg. you can find the minimum with a sort algorithm, but that would be dumb).

Sometimes informally used to describe problems where you can't even c

Stat <https://eduassistpro.github.io>

- Linear regression: $y_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik} + \epsilon$
- But only some of the covariates x_i subset gives the best MSE?
- 2^k possible subsets to check – increases exponentially in k .

$$2^{30} = 1,073,741,824$$

Require *approximate* solutions, often heuristic.

Some Caveats

Assignment Project Exam Help

- Big-O only relevant when n gets very large.

2

- <https://eduassistpro.github.io>

- You also care about readable understanda

Recursion, like the divine, can be pretty ineff

- Not only the number of operations matter, th
operation and the context make a differenc

Other Speed Considerations

■ Types of operation: multiplication/division take more time than addition/subtraction take more time than comparisons

■ <https://eduassistpro.github.io>

■ If **R** stores memory in RAM; if it runs out, it creates virtual RAM on your hard disk – this runs much slower

■ Small amount of memory within CPU is even faster (even for school).

■ Programming language also matters.

Compiled versus Interpreted Code

Most important distinction to be aware of.

- Operating systems provide the basic controls for computer hardware:

- Task scheduling.

-

-

<https://eduassistpro.github.io>

- Compiled code (C, C++, Fortran, COBOL) translated into a string of bit instructions that with the OS, before the code is executed.

- Interpreted code (R, Matlab, Java, Perl) translated into OS instructions as the program runs.

- Because of overhead in translating, interpreted code is *much* slower than compiled code.

Compiled versus Interpreted Code

So why interpreted code?

- Platform independent (if you have the right interpreter): R

- <https://eduassistpro.github.io>

- Fewer hassles (no memory allocation, eas sizes and types ...)

Many interpreted languages (including to be used to evaluate “chunks” of instructions muc

Many R built-in functions are pre-compiled.

Measuring Speed in R

Assignment Project Exam Help

R has some timing functions that can be useful to evaluate efficiency.

- <https://eduassistpro.github.io>

```
a = c() # vector of no length
system.time( for(i in 1:10000){ a = c(a, log
```

You can put a number of lines of code inside the call to `system.time` if you put everything inside `{ }`.

```
proc.time
```

`system.time` does exactly the following

```
start = proc.time() # starting time
```

```
nsim = 2
```

```
res = re
```

```
for(i in 1:nsim){
```

```
  X = rbinom(n,1,p)
```

```
  t[i] = sqrt(n)*abs( mean(X) - mu )/sd(X)
```

```
  res[i] = t[i] > qt(0.975,29)
```

```
}
```

```
proc.time()-start # time elapsed
```

Which I find easier to put down directly

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

```
proc.time()
```

Assignment Project Exam Help

```
> proc.time()
      user  system elapsed
    148      8
```

- <https://eduassistpro.github.io>
- Difference between the two is subtle and uni-
much time spent on user instructions versus
functions.
- `elapsed` is clock time; can vary depending on other processes
running.

R and Vectorization

R has compiled functions built in for vector/matrix operations

- `sum`, `mean`, `sd`, `var`, ...

- Matrix/vector multiplication and addition.

- Element-wise multiplication and addition and other built-in

These
used e

```
x = rnorm(100000)
```

```
start = proc.time()
```

```
m = x[1]
```

```
for(i in 2:length(x)){ m = ((i-1)/i)*m + x[i]/i }
```

```
proc.time()-start
```

```
system.time( {m2 = mean(x)} )
```

Making Use of Vectorization

Loops cannot always be avoided, but always ask “Could I do this with a vector?”

Eg: never loop through a vector if you are just doing arithmetical oper

Com



■ `sqrt`, `log`, `exp`, `dnorm`, ...

Vectorised operations:

■ `mean`, `sum`, `var`, `sd`, `cumsum`, `dif`

Matrix-vector operations:

■ `t`, `%*`, `%x`, `diag`, `solve`

Vectorization and Linear Algebra

Linear algebra often helps: taking column means of a matrix

```
X = matrix(rnorm(1000*5000),1000,5000)
```

```
ms = rep(0,ncol(X))
```

```
for(
```

Alter
mult

<https://eduassistpro.github.io>

Add WeChat [edu_assist_pro](#)

$$\begin{pmatrix} \frac{1}{n} \sum_{i=1}^n x_{i1} \\ \vdots \\ \frac{1}{n} \sum_{i=1}^n x_{ip} \end{pmatrix}$$

In code:

```
ms2 = rep(1/nrow(X),nrow(X))%*%X
```

But remember: clarity vs efficiency trade-off!

apply Functions

`apply` allows you to apply a function to the rows or columns of a matrix

```
ms3 = apply(x, 2, mean)
```

-
-
-

<https://eduassistpro.github.io>

Not actually any faster than a `for` loop

`lapply` breaks up output by `id`

of elements. (Output for elements with `index==2,...`)

`lapply/sapply` applies to each element in a list (eg vectors of different lengths), differ in output format.

Summary

- The way a task is computed can have a big impact on run-time.

- The way run-time scales with tasks can be very important (but not always).

-

- <https://eduassistpro.github.io>

- elegant and efficient.

- But number of computations not the only de

- In R, vectorization can have a dramatic impact on computational efficiency; most important thing to think about.

- Both complexity and vectorization can cost code readability – requires a balance, and good commenting.